

Neural Network Solvers for Combinatorial Optimization

Graph 11. Other CO Solvers (LLMs)

1

Outline of the CO Lectures

- Introduction to ML for Combinatorial Optimization (CO)
- Autoregressive (AR) CO Solvers
 - Reinforcement learning with RNN-based networks [Bello*, Pham*, et al., **ICLR 2017**]
 - Reinforcement learning with Transformer-based networks [W Kool et al., **ICLR 2019**]
- Non-autoregressive (NAR) CO Solvers
 - Reinforcement learning with DIMES [R Qiu*, Z Sun* & Y Yang, **NearIPS 2022**]
 - Supervised learning with DIFUSCO [Z Sun & Y Yang, **NearIPS 2023**]
- Pre-trained Large Language Models [C Yang et al., **ICLR 2024**]

2

Large Language Models as Optimizers,

(Google DeepMind: C Yang*, C Chen*, et al., ICLR 2024)

Key Idea



- Proposing **OPRO** (Optimization by Prompting) as a generic optimizer for solving **any problems** (e.g., regression or TSP) described in **natural language**;
- Each prompt contains a task description and a few solution/score pairs for previously solved problem instances;
- Using a LLM (with the prompt) to generate solutions for each new problem instance;
- Evaluate each new solution and adding the new solution/score pair to the prompt;
- Repeating the above steps until a termination condition is met.

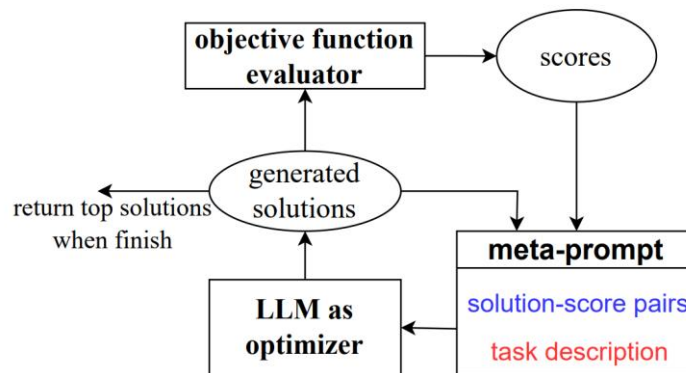
4/4/2024

@Yiming Yang, 11-741 S24 Graph11 Other CO Solvers (LLMs)

3

3

Learn to Solve Problems with OPRO



Yang, Chengrun, et al. "Large language models as optimizers." arXiv preprint arXiv:2309.03409 (2023).

4

4

Showcase of OPRO in Solving TSP

You are given a list of points with coordinates below: (0): (-4, 5), (1): (17, 76), (2): (-9, 0), (3): (-31, -86), (4): (53, -35), (5): (26, 91), (6): (65, -33), (7): (26, 86), (8): (-13, -70), (9): (13, 79), (10): (-73, -86), (11): (-45, 93), (12): (74, 24), (13): (67, -42), (14): (87, 51), (15): (83, 94), (16): (-7, 52), (17): (-89, 47), (18): (0, -38), (19): (61, 58).

Below are some previous traces and their lengths. The traces are arranged in descending order based on their lengths, where lower values are better.

<trace> 0,13,3,16,19,2,17,5,4,7,18,8,1,9,6,14,11,15,10,12 </trace>
length:
2254

<trace> 0,18,4,11,9,7,14,17,12,15,10,5,19,3,13,16,1,6,8,2 </trace>
length:
2017

<trace> 0,11,4,13,6,10,8,17,12,15,3,5,19,2,1,18,14,7,16,9 </trace>
length:
1953

<trace> 0,10,4,18,6,8,7,16,14,11,2,15,9,1,5,19,13,12,17,3 </trace>
length:
1840

Give me a new trace that is different from all traces above, and has a length lower than any of the above. The trace should traverse all points exactly once. The trace should start with <trace> and end with </trace>.

Input (Part I):
Node coordinates

Input (Part II):
Previous solutions

Instruction

Figure 18: An example of the meta-prompt for Traveling Salesman Problems with problem size $n = 20$. The blue text contains solution-score pairs; the orange text are meta-instructions.

5

Yang, Chengrun, et al. "Large language models as optimizers." arXiv preprint arXiv:2309.03409 (2023).

5

Evaluation Results

n	optimality gap (%)					# steps (# successes)		
	NN	FI	text-bison	gpt-3.5-turbo	gpt-4	text-bison	gpt-3.5-turbo	gpt-4
10	13.0 ± 1.3	3.2 ± 1.4	0.0 ± 0.0	0.0 ± 0.0	0.0 ± 0.0	40.4 ± 5.6 (5)	46.8 ± 9.3 (5)	9.6 ± 3.0 (5)
15	9.4 ± 3.7	1.2 ± 0.6	4.4 ± 1.3	1.2 ± 1.1	0.2 ± 0.2	N/A (0)	202.0 ± 41.1 (4)	58.5 ± 29.0 (4)
20	16.0 ± 3.9	0.2 ± 0.1	30.4 ± 10.6	4.4 ± 2.5	1.4 ± 0.6	N/A (0)	438.0 ± 0.0 (1)	195.5 ± 127.6 (2)
50	19.7 ± 3.1	9.8 ± 1.5	219.8 ± 13.7	133.0 ± 6.8	11.0 ± 2.6	N/A (0)	N/A (0)	N/A (0)

- Baseline NN (Nearest Neighbor Heuristic)
 - At each step, select the closest node from the current partial solution
- Baseline FI (Farthest Insertion)
 - At each step, add a new node that maximize the minimal insertion cost which is defined as

$$c(k) = \min_{i,j} d(i, k) + d(k, j) - d(i, j)$$

6

Concluding Remarks

■ Concept proving



- ORPO shows that LLMs with prompts in a loop can learn to optimize (mimicking gradient descent?)

■ Main limitations



- It cannot scale to large graphs due to the limited input length; similarly, it cannot handle a large training set of <solution, value> pairs.

■ Strong baselines are missing

- Comparison with DIMES and DIFUSCO on graphs with $n=10000$ nodes?
- Comparison with classic exact solvers?
- Comparison with LLMs for code generation?