

Deep Learning Techniques

DL 2. Recurrent Neural Networks (RNN)

1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

1

1

Outline

- Standard RNN (not Gated)
- Gated RNN


1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

2

2

RNN: Chained Building Blocks



An unrolled recurrent neural network.

1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

3

3

Consider a Vanilla Neural Network e.g., for hand-writing digit recognition

- Input layer**

$x \in \mathbb{R}^D$ is the feature vector of an image.

- Hidden layer**

$$h = g(W_{xh}x) = (g(x_1), g(x_2), \dots, g(x_D))$$

$$g(x_i) \stackrel{\text{def}}{=} \tanh(x_i) = \frac{e^{x_i} - e^{-x_i}}{e^{x_i} + e^{-x_i}} \quad (\text{rescaled logistic sigmoid})$$


- Output layer**

$$\hat{y} = f(\underbrace{W_{hy}h}_z) = (f_1(z), f_2(z), \dots, f_K(z))$$

$$f_j(z) = \text{softmax}(z) = \frac{\exp(z_j)}{\sum_{j'} \exp(z_{j'})}$$

the estimated probability of category j given $z = W_{hy}h$

Model Parameters $\theta = (W_{xh}, W_{hy})$



1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

4

4


Tanh function is a rescaled sigmoid

$$\sigma(x) \stackrel{\text{def}}{=} \frac{e^x}{1+e^x} \quad \text{for } x \in (-\infty, \infty),$$

$$\phi(x) = 2\sigma(2x) - 1$$

$$= 2 \frac{e^{2x}}{1+e^{2x}} - 1 = \frac{e^{2x}}{1+e^{2x}} = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

$$\stackrel{\text{def}}{=} \tanh(x)$$




Limitation of the Vanilla Neural Network

- **Cannot model** the input as **a sequence of tokens** but treating the input as a vector of in variables instead.
- **Cannot produce sequential output** but treating the output as a set of variables (with a probability score of each).
- **Cannot support language modeling**, the task to sequentially predict the next word based on the previous words

RNN for sequential modeling

(<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>)



An unrolled recurrent neural network.

1/23/2024

©Yiming Yang, S24 DL2 lecture on RNN

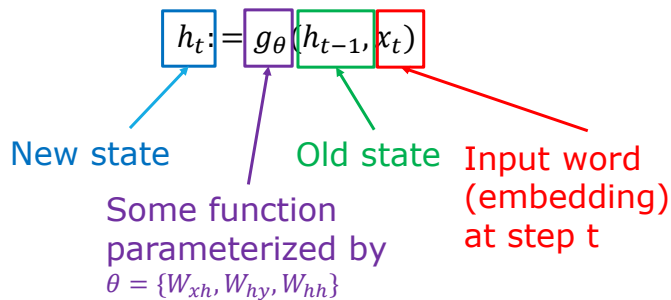
7

7


[Adapted from F.F. Li et al. Stanford CS231n]

Recurrent Neural Network (RNN)

- Modeling a sequence of (x_t, h_t, y_t) as



- The same θ used at each time step



1/23/2024

©Yiming Yang, S24 DL2 lecture on RNN

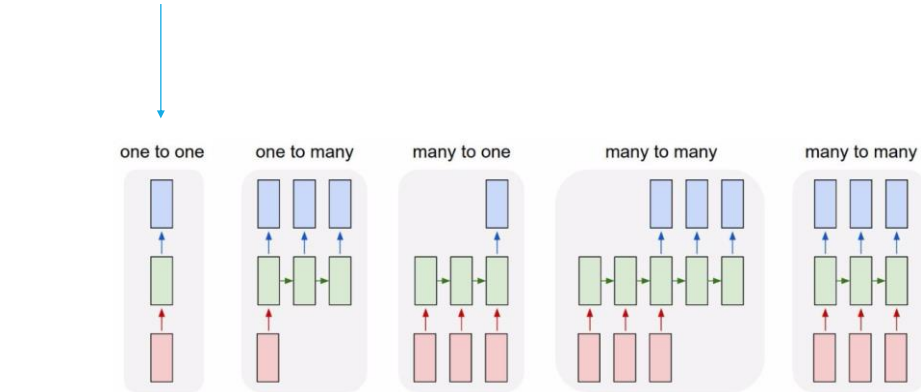
8

8

[Adapted from F.F. Li et al. Stanford CS231n]

Unrolled RNNs

- Vanila Neural Network (e.g., for image classification)



1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

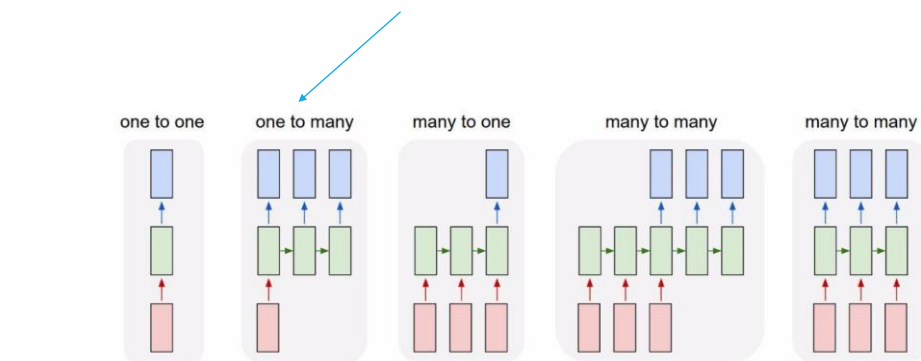
9

9

[Adapted from F.F. Li et al. Stanford CS231n]

Unrolled RNNs

E.g., Image Captioning: Image \rightarrow word sequence



1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

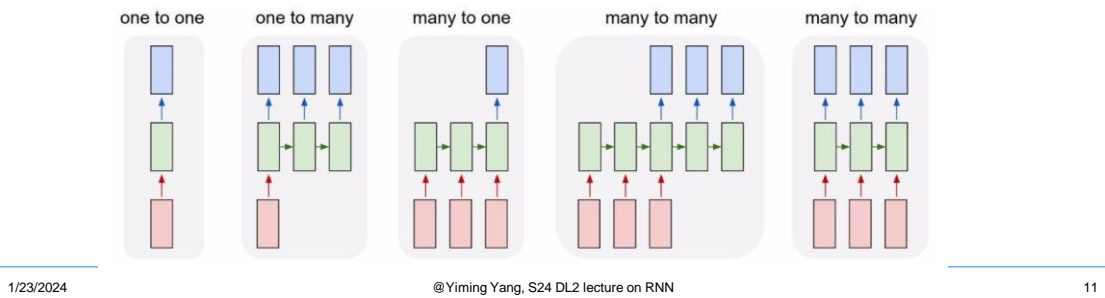
10

10

[Adapted from F.F. Li et al. Stanford CS231n]

Unrolled RNNs

E.g., Text classification:


Word sequence \rightarrow Topic label

11

[Adapted from F.F. Li et al. Stanford CS231n]

Unrolled RNNs

E.g., Machine Translation

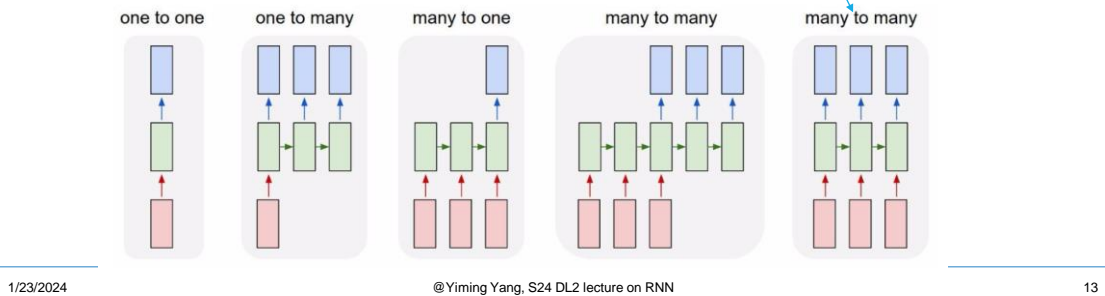
Word seq. \rightarrow Word seq.

12

[Adapted from F.F. Li et al. Stanford CS231n]

Unrolled RNNs

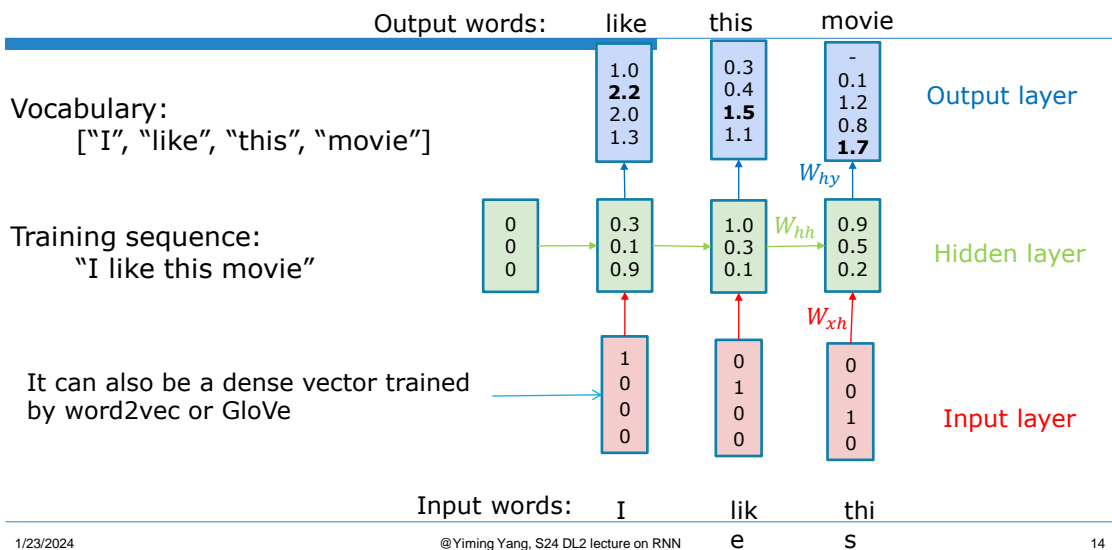
E.g., POS tagging
Word seq. \rightarrow Tag seq.



13

Example: Language Modeling

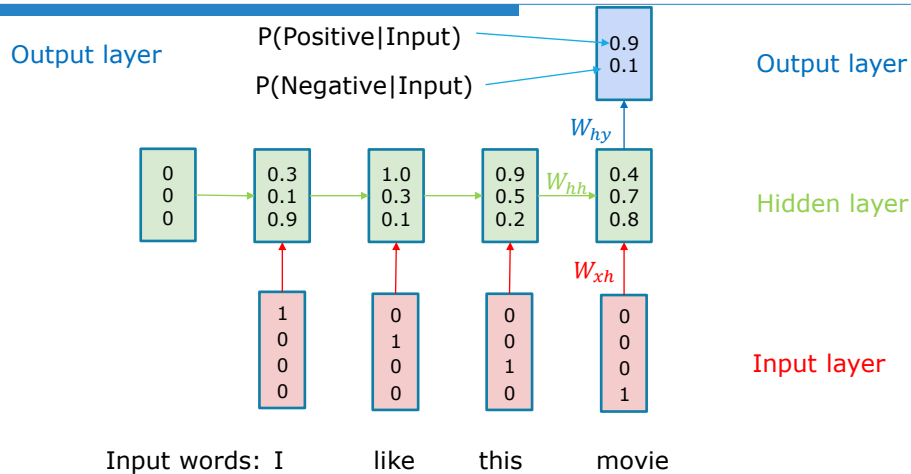
[Adapted from F.F. Li et al. Stanford CS231n]



14

[Adapted from F.F. Li et al. Stanford CS231n]

Example: Sentiment Classification



1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

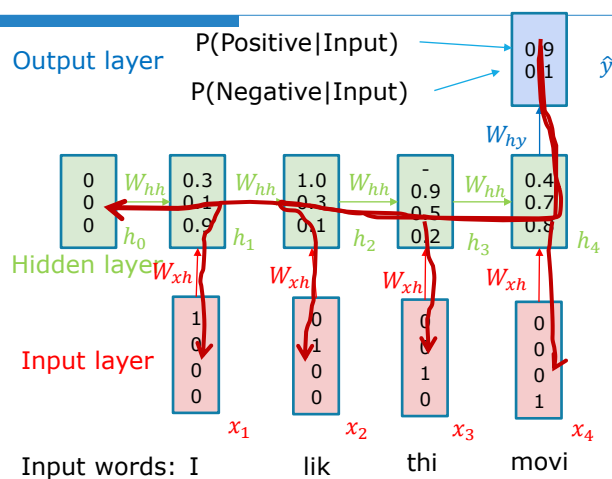
15

15

[Adapted from F.F. Li et al. Stanford CS231n]

Gradient vanishing problem in training RNN

- Multiplying the same matrices multiple times during forward/backwards propagation
- Causing the gradient to become very small or very large quickly, which is called the **gradient vanishing or exploding problem**



1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

16

16

Why do gradients vanish (or explode) in neural nets?

- “Neural Networks and Deep Learning”, Chapter 5, by Michael Nielsen (Dec 2017), <http://neuralnetworksanddeeplearning.com/chap5.html>
- Example: a multi-layer nnet with a sigmoid function at each layer
- Showing how the gradient of the output variable w.r.t. an input-layer variable would vanish or explode when the number of layers increases
- More generally, “neural networks suffer from an **unstable gradient** problem.”

1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

17

17

[based on the RNN tutorial by WILDML, 2015]

Why do gradients vanish in RNN?

- We could (incorrectly) compute the gradient of loss function $L(y, \hat{y})$ as


$$\frac{\partial L}{\partial W_{hh}} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial h_4} \frac{\partial h_4}{\partial W_{hh}}$$

- But h_4 also depends on h_3, h_2, h_1 , and each h_i depends on W_{hh} .

So, the correct formula should be

$$\begin{aligned} \frac{\partial L}{\partial W_{hh}} &= \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial h_4} \frac{\partial h_4}{\partial W_{hh}} + \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial h_4} \frac{\partial h_4}{\partial h_3} \frac{\partial h_3}{\partial W_{hh}} \\ &\quad + \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial h_4} \frac{\partial h_4}{\partial h_3} \frac{\partial h_3}{\partial h_2} \frac{\partial h_2}{\partial W_{hh}} + \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial h_4} \frac{\partial h_4}{\partial h_3} \frac{\partial h_3}{\partial h_2} \frac{\partial h_2}{\partial h_1} \frac{\partial h_1}{\partial W_{hh}} \\ &= \sum_{i=1}^n \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial h_n} \left(\prod_{j=i+1}^n \frac{\partial h_j}{\partial h_{j-1}} \right) \frac{\partial h_i}{\partial W_{hh}} \quad (\text{here } n = 4) \end{aligned}$$

- When n is large (e.g., 500), we have a long chain of $\prod_j \frac{\partial h_j}{\partial h_{j-1}}$.



1/23/2024


@Yiming Yang, S24 DL2 lecture on RNN

18

18

[based on the RNN tutorial by WILDML, 2015]

Why do gradients vanish in RNN (cont'd)



Notice that $\frac{\partial \tanh(v_j)}{\partial v_j}$ is between 0 and 1 and could be very near zero -- multiplying it several times would result in the vanishing gradient.

19

19

Outline

- Optimizing Neural Network
 - Stochastic Gradient Descent (Recap)
- Recurrent Neural Network (RNN)
 - Vanilla RNN
 - Gated RNN

20

Outline

- ✓ Standard RNN (not Gated)
- Gated RNN

1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

21

21

Gated Recurrent Neural Networks

- **LSTM** (Long Short Term Memory) as a representative model
- Introducing gates to plain RNN, to shorten the information flow and to avoid non-linear operations (like tanh) as needed
- Meditating the gradient vanishing issue in RNN effectively


1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN


22

22

Plain RNN vs. LSTM



The repeating module in a standard RNN contains a single layer.



The repeating module in an LSTM contains four interacting layers.


1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

23

23

LSTM Module



- Cell state c (a vector) allows the information to flow through a linear path (instead of non-linear activations) in many steps, and hence is **less prone to vanishing gradient**.
- Gates (σ 's) allow the system to skip some steps of the process, or to block the long chain of steps.

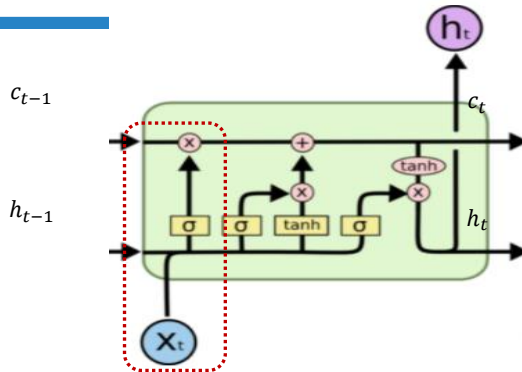
1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

24

24

LSTM Module



Forget gate $f_t = \sigma(W_f * [h_{t-1}, x_t] + b_f)$

Input gate $i_t = \sigma(W_i * [h_{t-1}, x_t] + b_i)$

Output gate $o_t = \sigma(W_o * [h_{t-1}, x_t] + b_o)$

The behavior of the gates are controlled by model parameters of W 's and b 's.

- $f_t = \sigma(\cdot) \in (0,1)$ mimic a soft gate.
- $f_t \rightarrow 0$ means the gate is closed, forcing c_{t-1} (old memory) to be forgotten.
- $f_t \rightarrow 1$ means the gate is open, allowing c_{t-1} (old memory) to be kept.


1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

25

25

LSTM Module



- Linear path

$$c_t = f_t * c_{t-1} + i_t * g_t$$

- Non-linear function

$$g_t = \tanh(W_c * [h_{t-1}, x_t] + b_c)$$

- If i_t is closed, then input x_t is skipped in the updating of c_t .
- If i_t is open but f_t is closed, then we forget c_{t-1} and renew c_t with g_t


1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

26

26

LSTM Module



Output gate $o_t = \sigma(W_o * [h_{t-1}, x_t] + b_o)$

$$h_t = o_t * \tanh(c_t)$$

- If o_t is closed, then we block the hidden state h_t from going forward.
- Otherwise, hidden state $h_t = \tanh(c_t)$ is propagated forward.

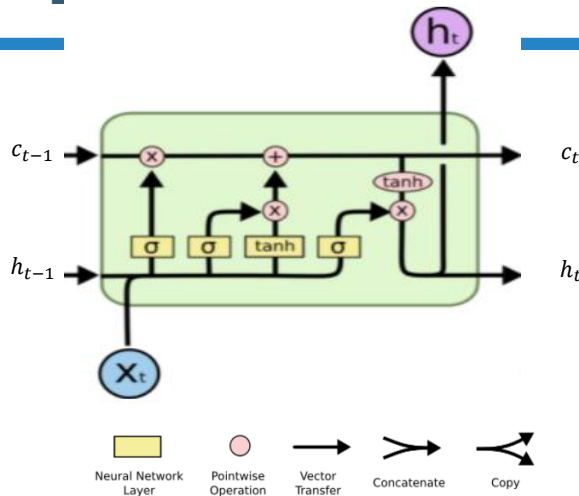
1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

27

27

LSTM Module



- Who control the gates?

- The model parameters W_f, W_i, W_o, W_c
- Not human designers of LSTM!

1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

28

28

Performance on language modeling

[Yoon Kim et al. AAAI 2016]

	PPL	Size
LSTM-Word-Small	97.6	5 m
LSTM-Char-Small	92.3	5 m
LSTM-Word-Large	85.4	20 m
LSTM-Char-Large	78.9	19 m
KN-5 (Mikolov et al. 2012)	141.2	2 m
RNN ¹ (Mikolov et al. 2012)	124.7	6 m
RNN-LDA ¹ (Mikolov et al. 2012)	113.7	7 m
genCNN ¹ (Wang et al. 2015)	116.4	8 m
FOFE-FNNLM ¹ (Zhang et al. 2015)	108.0	6 m
Deep RNN (Pascanu et al. 2013)	107.5	6 m
Sum-Prod Net ¹ (Cheng et al. 2014)	100.0	5 m
LSTM-1 ¹ (Zaremba et al. 2014)	82.7	20 m
LSTM-2 ¹ (Zaremba et al. 2014)	78.4	52 m

5-gram LM (not neural network)

Plain RNN for LM

Word-level LSTM

Table 3: Performance of our model versus other neural language models on the English Penn Treebank test set. *PPL* refers to perplexity (lower is better) and size refers to the approximate number of parameters in the model. KN-5 is a Kneser-Ney 5-gram language model which serves as a non-neural baseline. ¹For these models the authors did not explicitly state the number of parameters, and hence sizes shown here are estimates based on our understanding of their papers or private correspondence with the respective authors.

1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

29

29

Reference

- Hochreiter, Sepp, and Jürgen Schmidhuber. "Long short-term memory." *Neural computation* 9.8 (1997): 1735-1780.
- Louis-Philippe Morency, Tadas Baltrusaitis, **CMU 11-777: Advanced Multimodal Machine Learning**
- Fei-Fei Li, Andrej Karpathy, Justin Johnson [Stanford CS231n: Convolutional Neural Networks for Visual Recognition](#)
- RNN tutorial by WILDML <http://www.wildml.com/2015/10/recurrent-neural-networks-tutorial-part-3-backpropagation-through-time-and-vanishing-gradients/>
- [Christopher Olah's blog: Understanding LSTM Networks](#)
- [Denny Britz: Recurrent Neural Networks tutorial](#)

1/23/2024

@Yiming Yang, S24 DL2 lecture on RNN

30

30