

Classification

- CLS 3a. Stochastic Gradient Decent
- CLS 3b. Evaluation Metrics

2/13/2024

@Yiming Yang, lecture on SGD & EVAL

1

1

How fast does GE converge?

■ Theorem

If ℓ is both convex and differentiable ¹

$$\ell(w^{(k)}) - \ell(w^*) \leq \begin{cases} \frac{\|w^{(0)} - w^*\|_2^2}{2\eta k} = O\left(\frac{1}{k}\right) & \ell \text{ is convex} \\ \frac{c^k L \|w^{(0)} - w^*\|_2^2}{2} = O(c^k) & \ell \text{ is strongly convex} \end{cases} \quad (4)$$

where k is the number of iterations and $c \in (0, 1)$.

1: the step size η must be no larger than $\frac{1}{L}$ where L is the Lipschitz constant satisfying $\|\nabla \ell(a) - \nabla \ell(b)\|_2 \leq L\|a - b\|_2 \quad \forall a, b$

2/13/2024

@Yiming Yang, lecture on SGD & EVAL

5

5

How fast does GD converge?

Theorem

If ℓ is both convex and differentiable¹

$$\ell(w^{(k)}) - \ell(w^*) \leq \begin{cases} \frac{\|w^{(0)} - w^*\|_2^2}{2\eta^k} = O\left(\frac{1}{k}\right) & \ell \text{ is convex} \\ \frac{c^k L \|w^{(0)} - w^*\|_2^2}{2} = O(c^k) & \ell \text{ is strongly convex} \end{cases} \quad (4)$$

where k is the number of iterations and $c \in (0, 1)$.

- In general, to achieve $\ell(w^{(k)}) - \ell(w^*) \leq \rho$, GD needs $O\left(\frac{1}{\rho}\right)$ iterations;
- With strong convexity, it takes $O\left(\log\left(\frac{1}{\rho}\right)\right)$ iterations²

2: Convex Optimization, S. Boyd & L. Vandenberghe, Ch 9.3

2/13/2024

@Yiming Yang, lecture on SGD & EVAL

6

6

Why not happy with GD?

- Fast convergence \neq high efficiency

$$w^{(k)} = w^{(k-1)} - \eta_k \nabla \ell(w^{(k-1)}) \quad (5)$$

$$= w^{(k-1)} - \eta_k \nabla \left[\frac{1}{n} \sum_{i=1}^n \ell_i(w^{(k-1)}) \right] \quad (6)$$

- Per-iteration complexity = $O(n)$ (extremely large); a single iteration may take forever.
- How to make it cheaper? GD \Rightarrow SGD

2/13/2024

@Yiming Yang, lecture on SGD & EVAL

7

7

Stochastic Gradient Decent

Approximate the full gradient via an unbiased estimator

$$w^{(k)} = w^{(k-1)} - \eta_k \nabla \left(\frac{1}{n} \sum_{i=1}^n \ell_i(w^{(k-1)}) \right) \quad (7)$$

$$\approx \underbrace{w^{(k-1)} - \eta_k \nabla \left(\frac{1}{|B|} \sum_{i \in B} \ell_i(w^{(k-1)}) \right)}_{\text{mini-batch SGD}^3} \quad B \stackrel{\text{unif}}{\sim} \{1, 2, \dots, n\} \quad (8)$$

$$\approx \underbrace{w^{(k-1)} - \eta_k \nabla \ell_i(w^{(k-1)})}_{\text{pure SGD}} \quad i \stackrel{\text{unif}}{\sim} \{1, 2, \dots, n\} \quad (9)$$

Trade-off: lower computation cost v.s. larger variance

³When using GPU, $|B|$ usually depends on the memory budget.

GD vs SGD

For strongly convex $\ell(w)$, according to [Bottou, 2012]

Optimizer	GD	SGD	Winner
Time per-iteration	$O(n)$	$O(1)$	SGD
Iterations for accuracy ρ	$O\left(\log\left(\frac{1}{\rho}\right)\right)$	$\tilde{O}\left(\frac{1}{\rho}\right)$	GD
Time for accuracy ρ	$O\left(n \log \frac{1}{\rho}\right)$	$\tilde{O}\left(\frac{1}{\rho}\right)$	Depends
Time for test-set error ϵ	$O\left(\frac{1}{\epsilon^{1/\alpha}} \log \frac{1}{\epsilon}\right)$	$\tilde{O}\left(\frac{1}{\epsilon}\right)$	SGD

where $\frac{1}{2} \leq \alpha \leq 1$

SVM Solver: Pegasos [Shalev-Shwartz et al., 2011]

$$\ell_i(w) = \max(0, 1 - y_i w^\top x_i) + \frac{\lambda}{2} \|w\|^2 \quad (10)$$

$$= \begin{cases} \frac{\lambda}{2} \|w\|^2 & y_i w^\top x_i \geq 1 \\ 1 - y_i w^\top x_i + \frac{\lambda}{2} \|w\|^2 & y_i w^\top x_i < 1 \end{cases} \quad (11)$$

Therefore

$$\nabla \ell_i(w) = \begin{cases} \lambda w & y_i w^\top x_i \geq 1 \\ \lambda w - y_i x_i & y_i w^\top x_i < 1 \end{cases} \quad (12)$$

2/13/2024

@Yiming Yang, lecture on SGD & EVAL

10

10

SVM Solver in 10 Lines

Algorithm 1: Pegasos: SGD solver for SVMs

Input: n, λ, T ;

Initialization: $w \leftarrow 0$;

for $k = 1, 2, \dots, T$ **do**

$i \stackrel{\text{uni}}{\sim} \{1, 2, \dots, n\}$;

$\eta_k \leftarrow \frac{1}{\lambda k}$;

if $y_i w^{(k)\top} x_i < 1$ **then**

$w^{(k+1)} \leftarrow w^{(k)} - \eta_k (\lambda w^{(k)} - y_i x_i)$

else

$w^{(k+1)} \leftarrow w^{(k)} - \eta_k \lambda w^{(k)}$

end

end

Output: $w^{(T+1)}$

2/13/2024

@Yiming Yang, lecture on SGD & EVAL

11

11

Empirical Comparison

SGD v.s. batch solvers⁴ on RCV1

#Features	#Training examples
47,152	781,265

Algorithm	Time (secs)	Test Error
SMO (SVM ^{light})	≈ 16,000	6.02%
Cutting Plane (SVM ^{perf})	≈ 45	6.02%
SGD	< 1	6.02%

What is the magic?


2/13/2024

⁴<http://leon.bottou.org/projects/sgd>

12

12

The Magic



- ▶ SGD takes a long time to reach an highly accurate solution on the training data.
- ▶ However, for low generalization error (on test data) we may not need the training-set accuracy to be too high.

2/13/2024


⁴<http://leon.bottou.org/slides/largescale/1stut.pdf>

13

13

Empirical Results

On Netflix [Gemulla et al., 2011]



DSGD, PSGD: Distributed SGD

ALS: Alternating least squares, one of the state-of-the-art batch solvers

DGD: Distributed GD

2/13/2024

14

14

Popular SGD Variants

A non-exhaustive list

1. AdaGrad [Duchi et al., 2011]
2. Momentum [Rumelhart et al., 1988]
3. Nesterov's method [Nesterov et al., 1994]
4. AdaDelta: AdaGrad refined [Zeiler, 2012]
5. Rprop & Rmsprop [Tieleman and Hinton, 2012]
6. Stochastic Variance Reduced Gradient [Johnson and Zhang, 2013]
7. ADAM [Kingma and Ba, 2014]

All are empirically found effective in solving nonconvex problems (e.g., deep neural nets).

Demos ⁶: Animation 0, 1, 2, 3

2/13/2024

15

⁶https://www.reddit.com/r/MachineLearning/comments/2gopfa/visualizing_gradient_optimization_techniques/cklhott

15

Summarization Remarks

- SGD is extremely handy for large problems.
- It's only one of many handy tools
 - Alternatives: quasi-Newton (BFGS), Coordinate descent, ADMM, CG, etc.
 - How to choose? Depending on the problem structures.

Classification

- CLS 3a. Stochastic Gradient Decent
- CLS 3b. Evaluation Metrics

Testing-phase Decision Making

- **Binary Classification**

- Using one binary classifier to make a yes/no decision for each test instance.

- **Multi-label Classification**

- Using K OVA (one-vs-all) classifiers to make an independent yes/no decision per category for each test instance.

- **Multi-class Classification**

- Using a multi-class model (e.g., softmax or kNN) to pick one label (out of K) for each test instance.

2/13/2024

@Yiming Yang, lecture on SGD & EVAL

20

20

Binary Classification Evaluation (e.g., sentiment classification)

- Given a test set of n instances, the system make n yes/no decisions, as summarized by a **two-by-two contingency table**

	$y = 1$	$y = 0$	
$\hat{y} = 1$	a	b	$a + b$
$\hat{y} = 0$	c	d	$c + d$
	$a + c$	$b + d$	$n = a + b + c + d$

- Here a (true positive), b (false positive), c (false negative) and d (true negative) are the counts of test instances in the four categories

2/13/2024

@Yiming Yang, lecture on SGD & EVAL

21

21

Evaluation Metrics for Binary Classifier

- Precision $p = \frac{a}{a+b}$
- Recall $r = \frac{a}{a+c}$
- F_1 -measure $F_1 = \frac{2pr}{p+r} = \frac{2a}{2a+b+c}$
- Accuracy $Acc = \frac{a+d}{n}$
- Error $Err = \frac{b+c}{n}$

Why do we prefer F_1 over p and r ?

- A trivial classifier always predicting "yes" for every input will have $r = 100\%$ but is totally useless.
- A trivial classifier always predicting the most popular label without checking the input is totally useless, but would have a high precision, especially when the labels are unbalanced in the data collection.
- F_1 is the **harmonic average** of p and r , which means that the F_1 value is high only if p and r are both high.

Harmonic Average

Def. $f(x_1, \dots, x_n) = \frac{1}{\frac{1}{n}(\frac{1}{x_1} + \dots + \frac{1}{x_n})}$

$$F_1(p, r) = \frac{2pr}{r+p} = \frac{1}{\frac{1}{2}(\frac{1}{p} + \frac{1}{r})}$$

← Harmonic average of p and r
← Dominated by the smaller one

$$F_\beta(p, r) = \frac{(1+\beta^2)pr}{r+\beta^2p} = \frac{1}{\frac{1}{1+\beta^2}\frac{1}{p} + \frac{\beta^2}{1+\beta^2}\frac{1}{r}}$$

← Weighted harmonic average

2/13/2024


@Yiming Yang, lecture on SGD & EVAL

24


24

F_β with different values of β

$$F_\beta(p, r) = \frac{1}{\frac{1}{1+\beta^2}\frac{1}{p} + \frac{\beta^2}{1+\beta^2}\frac{1}{r}}$$



$$F_{\beta=0} = \frac{1}{\frac{1}{1+0}\frac{1}{p} + \frac{0}{1+0}\frac{1}{r}} = p$$



$$F_{\beta=\infty} = \frac{1}{\frac{1}{1+\infty}\frac{1}{p} + \frac{\infty}{1+\infty}\frac{1}{r}} = r$$

Does $F_{\beta=0.5}$ favor r or p ? What about $F_{\beta=2}$?

2/13/2024

@Yiming Yang, lecture on SGD & EVAL

25

25

Multi-class Classification Evaluation

(e.g., OCR word recognition, Amazon product ID, etc.)

- The system assign one and only one label per test instance, which can be summarized using a K -by- K confusion matrix M .
- A toy example about 3-way movie classification to the right.
- $Acc = \frac{1}{n} \sum_{k=1}^K M_{kk} = \frac{1+76+1}{100} = 78\%$
- $Err = 1 - Acc = 22\%$

	$y = \text{good}$	$y = \text{soso}$	$y = \text{bad}$	
$\hat{y} = \text{good}$	1	3	1	5
$\hat{y} = \text{soso}$	20	70	0	90
$\hat{y} = \text{bad}$	3	1	1	5

Confusion Matrix

2/13/2024

@Yiming Yang, lecture on SGD & EVAL

26

26

Multi-label Classification Evaluation

(e.g., wiki page classification by an OVA model per label)

- For $k = 1, 2, \dots, K$, construct a contingency table per category as

	$y_k = 1$	$y_k = 0$	
$\hat{y}_k = 1$	a_k	b_k	$a_k + b_k + c_k + d_k = n$
$\hat{y}_k = 0$	c_k	d_k	

- Micro-averaging
 - Merge the K tables into one table by summing the corresponding cells
 - Use the merged table to compute p, r and F_1 (but not Acc or Err)
- Macro-averaging
 - Use each of the K tables to compute the category-specific p, r and F_1 , then average them over categories as $\bar{p} = \frac{1}{K} \sum_{k=1}^K p_k$, and so on.

2/13/2024

@Yiming Yang, lecture on SGD & EVAL

27

27

Micro-averaging vs. Macro-averaging

- Micro-averaging gives each test instance an equal weight
- Macro-average gives each category an equal weight
- Both are informative for method comparison

2/13/2024

@Yiming Yang, lecture on SGD & EVAL

28

28

Issues with Accuracy or Error (especially for rare categories)

- Consider when the number of labels is large, e.g., $K = 1000$.
- On average, 99.9% of the documents are negative instances for each OVA model.
- Consider a trivial classifier Mr. NO, with the contingency table like

	$y_k = 1$	$y_k = 0$	
$\hat{y}_k = 1$	0	0	$Acc = \frac{a_k + d_k}{n} = 99.9\%$ $Err = \frac{b_k + c_k}{n} = 0.1\%$
$\hat{y}_k = 0$	c_k	d_k	

- Take-home message: Focus on F_1 (or F_β) instead.

2/13/2024

@Yiming Yang, lecture on SGD & EVAL

29

29